

UseCase_Module 3:Index & Statistics Tuning

DAY 2 – Indexing, Statistics & Concurrency

Module 3: Index & Statistics Tuning (4 Hours)

Concepts

- Clustered vs Non-clustered indexes
- Covering & filtered indexes
- Columnstore indexes (when to use)
- Index fragmentation myths
- Fill factor strategy
- Statistics:
 - Histograms
 - Auto-update thresholds
 - Multi-column stats

Hands-on Labs

- Create optimal indexes for slow queries
- Analyze missing index DMVs
- Rebuild vs reorganize decisions
- Fix queries using statistics updates

Tools Used

- `sys.dm_db_index_usage_stats`
- `sys.dm_db_missing_index_details`
- `DBCC SHOW_STATISTICS`

Create CSC Global Use Cases for DBA professionals. Also explain each line of code.

Module 3: Index & Statistics Tuning (4 Hours)

I've structured it exactly how **senior DBA training** should run:

- Real **enterprise use cases**
 - **Hands-on labs**
 - **Every line of code explained**
 - Clear **why this matters in production**
-

DAY 2 – Indexing, Statistics & Concurrency

Module 3: Index & Statistics Tuning (4 Hours)

USE CASE 1 – “Overdue Compliance Report is CPU Heavy”

(Clustered vs Nonclustered + Covering Index)

Business Context (CSC Style)

The global compliance dashboard is slow every morning.

DBA suspects **scans + key lookups**.

Lab 1A – Observe the problem

```
SET STATISTICS IO ON;
SET STATISTICS TIME ON;

SELECT
    e.country_code,
    e.client_name,
    cf.filing_type,
    COUNT_BIG(*) AS overdue_count
FROM dbo.ComplianceFilings cf
JOIN dbo.Entities e
    ON e.entity_id = cf.entity_id
WHERE
    cf.filing_status = 'OVERDUE'
    AND cf.due_date < GETDATE()
GROUP BY
    e.country_code,
    e.client_name,
    cf.filing_type;
```

Line-by-line explanation

- SET STATISTICS IO ON
 - Shows logical reads → **true performance cost**
- SET STATISTICS TIME ON
 - Shows CPU & elapsed time
- JOIN dbo.Entities e ON e.entity_id = cf.entity_id
 - Join depends on index availability
- COUNT_BIG(*)

Safe aggregation for large datasets

- GROUP BY

Forces aggregation → may cause hash aggregate

Key DBA Learning

 **Scans + Aggregation = expensive if indexes are wrong**

Lab 1B – Create a covering nonclustered index

```
CREATE INDEX IX_CF_Status_DueDate_Cover
ON dbo.ComplianceFilings (filing_status, due_date)
INCLUDE (entity_id, filing_type);
```

Explanation

- filing_status, due_date → WHERE clause
- INCLUDE (entity_id, filing_type) → JOIN + GROUP BY columns
- This **eliminates key lookups**

Learning

 **Covering indexes reduce IO more than almost any other tuning**

USE CASE 2 – “DBA Added 50 Indexes... Performance Got Worse”

(Index Usage + Fragmentation Myths)

Lab 2A – Find unused indexes

```
SELECT
    OBJECT_NAME(i.object_id) AS table_name,
    i.name AS index_name,
    s.user_seeks,
    s.user_scans,
    s.user_lookups,
    s.user_updates
FROM sys.indexes i
LEFT JOIN sys.dm_db_index_usage_stats s
    ON i.object_id = s.object_id
```

```
AND i.index_id = s.index_id
WHERE OBJECTPROPERTY(i.object_id,'IsUserTable') = 1
ORDER BY s.user_seeks + s.user_scans + s.user_lookups;
```

Line-by-line explanation

- sys.indexes → all indexes
- sys.dm_db_index_usage_stats → how indexes are used
- user_seeks → good usage
- user_scans → may indicate poor design
- user_updates → maintenance overhead
- LEFT JOIN → includes unused indexes

Key DBA Learning

⚠️ Unused indexes slow INSERT/UPDATE/DELETE & waste memory

Index Fragmentation Myth (Important Teaching Point)

Fragmentation ≠ performance problem for OLTP seeks
Fragmentation mostly affects **range scans**

Lab 2B – Decide rebuild vs reorganize

```
SELECT
    OBJECT_NAME(object_id) AS table_name,
    index_id,
    avg_fragmentation_in_percent,
    page_count
FROM sys.dm_db_index_physical_stats
(
    DB_ID(),
    NULL,
    NULL,
    NULL,
    'SAMPLED'
)
WHERE page_count > 1000;
```

Explanation

- avg_fragmentation_in_percent
 - < 5% → ignore
 - 5–30% → reorganize
 - 30% → rebuild

- page_count > 1000

Ignore tiny indexes

Learning

 **Fragmentation decisions must be data-driven, not scheduled blindly**

USE CASE 3 – “Client-wise Reports Are Inconsistent”

(Statistics, Histograms & Cardinality)

Lab 3A – Inspect statistics histogram

```
DBCC SHOW_STATISTICS ('dbo.ComplianceFilings', 'IX_CF_Status_DueDate_Cover')
WITH HISTOGRAM;
```

Explanation

- DBCC SHOW_STATISTICS → internal optimizer data
- HISTOGRAM → distribution of column values
- Shows:
 - RANGE_HI_KEY
 - EQ_ROWS
 - RANGE_ROWS

Teaching Point

- SQL Server estimates rows **based on histogram steps**
- Skewed data = wrong estimates

Key DBA Learning

 **Bad statistics = bad plans even with good indexes**

Lab 3B – Show auto-update threshold problem

```
SELECT
    name,
    auto_created,
    STATS_DATE(object_id, stats_id) AS last_updated
FROM sys.stats
```

```
WHERE object_id = OBJECT_ID('dbo.ComplianceFilings');
```

Explanation

- auto_created = stats created automatically
- STATS_DATE = last refresh time

Teaching Point

- Auto-update happens only after ~20% + 500 rows change
- Large tables wait too long

Lab 3C – Fix with manual stats update

```
UPDATE STATISTICS dbo.ComplianceFilings WITH FULLSCAN;
```

Explanation

- FULLSCAN reads all rows
- Best estimates (slower but accurate)

Learning

 **Critical reporting tables need proactive stats maintenance**

USE CASE 4 – “Query Uses Multiple Columns but Estimates Are Wrong”

(Multi-column Statistics)

Lab 4A – Create multi-column statistics

```
CREATE STATISTICS ST_CF_Status_DueDate  
ON dbo.ComplianceFilings (filing_status, due_date);
```

Explanation

- Helps optimizer understand **column correlation**
- Single-column stats assume independence (often wrong)

Learning

 **Multi-column stats can fix estimates without adding indexes**

USE CASE 5 – “Missing Index DMV Lies”

(When & How to Use Missing Index DMVs)

Lab 5A – View missing index recommendations

```
SELECT
    migs.avg_total_user_cost,
    migs.avg_user_impact,
    mid.statement,
    mid.equality_columns,
    mid.inequality_columns,
    mid.included_columns
FROM sys.dm_db_missing_index_details mid
JOIN sys.dm_db_missing_index_groups mig
    ON mid.index_handle = mig.index_handle
JOIN sys.dm_db_missing_index_group_stats migs
    ON mig.index_group_handle = migs.group_handle
ORDER BY migs.avg_user_impact DESC;
```

Explanation

- avg_user_impact → estimated improvement %
- equality_columns → WHERE col = ?
- inequality_columns → WHERE col <, >
- included_columns → SELECT list

Teaching Warning

Never auto-create missing indexes

- They:
 - Ignore maintenance cost
 - Ignore existing indexes
 - Often overlap

Learning

Missing index DMV is a suggestion engine, not a solution engine

USE CASE 6 – Columnstore Index (When to Use)

Business Context

Compliance analytics, monthly reporting, read-heavy workload.

Lab 6A – Create Nonclustered Columnstore

```
CREATE NONCLUSTERED COLUMNSTORE INDEX NCCI_CF
ON dbo.ComplianceFilings
(
    filing_status,
    due_date,
    entity_id,
    filing_type
);
```

Explanation

- Columnstore compresses data
- Batch-mode execution
- Best for:
 - Large scans
 - Aggregations
 - Reporting

When NOT to use

- High-frequency OLTP updates
- Small tables

Learning

 **Columnstore is a reporting accelerator, not an OLTP index**

FINAL DAY-2 DBA TAKEAWAYS (CSC Level)

After this module, DBAs must know:

1 Indexes are **workload-specific**, not generic

-
- 2 Covering indexes reduce IO more than CPU tuning
 - 3 Statistics quality matters as much as indexes
 - 4 Fragmentation is often over-tuned
 - 5 Missing index DMVs must be interpreted, not trusted blindly
 - 6 Columnstore is powerful—but only in the right workload
-

Interview-Ready One-Liners (Use These)

- “Indexes shape plans; statistics shape estimates.”
 - “Fragmentation hurts scans, not seeks.”
 - “Missing index DMVs show symptoms, not design.”
 - “Bad stats can break even perfect indexes.”
-

DAY-2 DEMO PACK

(Bad Statistics → Bad Plan → Fix → Verify)

Goal for participants:

Understand how bad or stale statistics alone (even with indexes present) can produce wrong plans, high memory grants, scans, and slow performance—and how a DBA fixes it correctly.

DEMO STORY (CSC Global Context)

“Nightly compliance load added ~80,000 new filings.
Indexes exist, but the morning report is suddenly slow.”

Root cause: **Statistics did not refresh → optimizer underestimates rows → wrong plan.**

STEP 0 – Instructor Setup

-- Always enable these for demo clarity
SET NOCOUNT ON;
SET STATISTICS IO ON;

```
SET STATISTICS TIME ON;
```

-- Ask participants to enable:
-- Actual Execution Plan (Ctrl + M in SSMS)

 **Tell class:**

“Today, we are NOT fixing indexes first. We’ll prove stats alone can break plans.”

STEP 1 – Confirm Index Exists (Important!)

```
SELECT
    i.name AS index_name,
    i.type_desc
FROM sys.indexes i
WHERE i.object_id = OBJECT_ID('dbo.ComplianceFilings')
    AND i.name = 'IX_Filings_Status_DueDate';
```

Output Interpretation

- Index **exists**
- So if performance is bad → **indexes are NOT the problem**

 **Key Learning:**

If an index exists and query is still slow, look at statistics next.

STEP 2 – SIMULATE BAD STATISTICS

Force statistics to be stale

-- Simulate stale stats by disabling auto update (DEMO ONLY)

```
ALTER DATABASE CURRENT SET AUTO_UPDATE_STATISTICS OFF;
GO
```

 **Instructor note:**

Emphasize this is **for demo only**, not recommended blindly in prod.

STEP 3 – RUN QUERY WITH BAD STATS (BAD PLAN)

```
DECLARE @AsOfDate DATE = CAST(GETDATE() AS DATE);

SELECT
    e.country_code,
    e.client_name,
    cf.filing_type,
    COUNT_BIG(*) AS overdue_count
FROM dbo.ComplianceFilings cf
JOIN dbo.Entities e
    ON e.entity_id = cf.entity_id
WHERE
    cf.filing_status = 'OVERDUE'
    AND cf.due_date < @AsOfDate
GROUP BY
    e.country_code,
    e.client_name,
    cf.filing_type
ORDER BY
    overdue_count DESC;
GO
```

STEP 4 – WHAT DBAs SEE (BAD PLAN OUTPUT)

In Actual Execution Plan:

- **✗ Index Scan** instead of Seek
- **✗ Hash Match (Aggregate)** with large input
- **✗ Sort** operator
- **⚠ Possible spill warning**
- **⚠ Huge memory grant**

In Messages tab:

- High **logical reads**
- Higher **CPU & elapsed time**

 **Key DBA Learning:**

Bad statistics cause bad estimates → bad join + aggregate choices.

STEP 5 – PROVE STATS ARE WRONG

```
DBCC SHOW_STATISTICS
```

```

(
  'dbo.ComplianceFilings',
  'IX_Filings_Status_DueDate'
)
WITH HISTOGRAM;

```

Explain Output (Important Columns)

Column	Meaning
RANGE_HI_KEY	Upper bound of value range
EQ_ROWS	Exact matches
RANGE_ROWS	Rows between ranges

🎯 Instructor talking point:

“If the histogram says ~100 rows but actual is 50,000 → optimizer is blind.”

STEP 6 – FIX: UPDATE STATISTICS

UPDATE STATISTICS dbo.ComplianceFilings WITH FULLSCAN;

UPDATE STATISTICS dbo.Entities WITH FULLSCAN;
GO

Why FULLSCAN?

- Reads all rows
- Best histogram accuracy
- Perfect for reporting tables

🔑 Key DBA Learning:

Statistics quality is often more important than index quantity.

STEP 7 – VERIFY (SAME QUERY, BETTER PLAN)

DECLARE @AsOfDate2 DATE = CAST(GETDATE() AS DATE);

```

SELECT
  e.country_code,
  e.client_name,

```

```
cf.filing_type,  
COUNT_BIG(*) AS overdue_count  
FROM dbo.ComplianceFilings cf  
JOIN dbo.Entities e  
    ON e.entity_id = cf.entity_id  
WHERE  
    cf.filing_status = 'OVERDUE'  
    AND cf.due_date < @AsOfDate2  
GROUP BY  
    e.country_code,  
    e.client_name,  
    cf.filing_type  
ORDER BY  
    overdue_count DESC;  
GO
```

What improves

- Index **Seek**
 - Smaller hash/sort
 - Reduced memory grant
 - Lower reads
 - Faster runtime
-

STEP 8 – CLEANUP (IMPORTANT)

```
ALTER DATABASE CURRENT SET AUTO_UPDATE_STATISTICS ON;  
GO
```

```
SET STATISTICS IO OFF;  
SET STATISTICS TIME OFF;
```

🎯 DAY-2 CORE DBA TAKEAWAY

“Indexes decide access paths.
Statistics decide how many rows SQL Server thinks will flow.
Bad stats break good indexes.”

 **INDEX MAINTENANCE DECISION FLOWCHART**

(Interview + Production Ready)

START

|

v

Is the query slow?

|

++ NO --> Do nothing

|

v

Is index being used?

|

++ NO --> Check missing index DMV

|

v

Is index heavily updated?

|

++ YES --> Check index usage stats

| |

++ user_updates high & seeks low?

| |

++ YES --> Drop or redesign index

|

v

Check fragmentation (page_count > 1000)

|

++ < 5% --> Ignore

|

++ 5–30% --> REORGANIZE

|

++ >30% --> REBUILD

|

v

Is query still slow?

|

++ YES --> Check statistics freshness

| |

++ STATS_DATE old?

| |

++ YES --> UPDATE STATISTICS

|

v

Are estimates still wrong?

|

++ YES --> Create multi-column stats

```
|  
v  
Is workload analytical (large scans)?  
|  
++- YES --> Consider COLUMNSTORE  
|  
v  
END (Query stable)
```

Interview-Ready DBA One-Liners

- “Fragmentation hurts scans, not seeks.”
 - “Never rebuild indexes blindly—check page count.”
 - “Missing index DMV is a suggestion, not a command.”
 - “Bad statistics break good indexes.”
 - “Columnstore is for analytics, not OLTP.”
-

4 CSC Global-style DBA use cases for Module 3 (Index & Statistics Tuning).

Each use case includes:

- **Scenario (what CSC DBAs face)**
 - **Problem symptoms**
 - **Bad query / missing support**
 - **Step-by-step code with line-by-line explanation**
 - **What “good plan” should look like**
 - **How to validate via DMVs / SHOW_STATISTICS**
-

Use Case 1 — Covering Index to Remove Key Lookups (High logical reads)

Scenario

CSC portal shows “Open invoices for an entity”. Query is fast sometimes, slow for entities with many invoices because of **Key Lookups** (bookmark lookups) on a large table.

Bad Query (works but inefficient)

```
SELECT invoice_id, due_date, amount, currency_code  
FROM dbo.Invoices  
WHERE entity_id = @EntityId  
AND invoice_status = 'Open'
```

```
ORDER BY due_date DESC;
```

Why it becomes slow

- Index may exist on (entity_id, invoice_status) but selected columns aren't covered → SQL Server does **Key Lookup** for each matching row.
 - Many rows = many lookups = high reads.
-

Step-by-step Fix (Covering index)

Code

```
CREATE INDEX IX_Invoices_EntityStatus_DueDate_Cover
ON dbo.Invoices (entity_id, invoice_status, due_date DESC)
INCLUDE (amount, currency_code);
GO
```

Line-by-line explanation

- **CREATE INDEX IX_Invoices_EntityStatus_DueDate_Cover**
Creates a new **nonclustered index** with a meaningful name.
- **ON dbo.Invoices (entity_id, invoice_status, due_date DESC)**
Puts the most selective filters first (entity_id, invoice_status).
Adds due_date DESC to support the ORDER BY without extra sort.
- **INCLUDE (amount, currency_code)**
Stores these columns in the index leaf level so SQL Server can return them **without key lookups**.
- **GO**
Ends the batch.

Expected plan improvement

- **Index Seek** on this new index
 - **No Key Lookup**
 - Often **no Sort operator**
-

Validate improvement (usage tracking)

```
SELECT *
FROM sys.dm_db_index_usage_stats
WHERE database_id = DB_ID()
AND object_id = OBJECT_ID('dbo.Invoices');
```

Explanation

- Shows how often each index is used:
 - user_seeks (good)
 - user_scans
 - user_lookups
 - user_updates (write cost)

Use Case 2 — Filtered Index for Rare Status (Overdue Filings)

Scenario

CSC compliance team frequently checks **Overdue filings** (rare data).

A full index on status may still be heavy because the table is huge.

Query

```
SELECT TOP (2000) filing_id, entity_id, due_date, penalty_amount
FROM dbo.ComplianceFilings
WHERE filing_status = 'Overdue'
ORDER BY due_date DESC;
```

Fix: Filtered Index (targeted & small)

Code

```
CREATE INDEX IX_Filings_Overdue_DueDate
ON dbo.ComplianceFilings (due_date DESC)
INCLUDE (entity_id, penalty_amount)
WHERE filing_status = 'Overdue';
GO
```

Line-by-line explanation

- CREATE INDEX IX_Filings_Overdue_DueDate

Creates a dedicated index specifically for “Overdue” status queries.

- ON dbo.ComplianceFilings (due_date DESC)

Sort key supports ORDER BY due_date DESC.

- INCLUDE (entity_id, penalty_amount)

Covers returned columns; avoids lookups.

- WHERE filing_status = 'Overdue'

This is the **filtered condition**.

Only “Overdue” rows are stored in this index → far smaller → faster seeks.

Expected plan improvement

- Seek on filtered index
- No scan on main status index
- Very low logical reads

Validate filtered index suggestion impact (Missing index DMV)

```
SELECT TOP (20)
    mid.statement,
    mid.equality_columns,
    mid.inequality_columns,
    mid.included_columns
FROM sys.dm_db_missing_index_details AS mid
ORDER BY mid.index_handle DESC;
```

Explanation

- sys.dm_db_missing_index_details shows **potential** missing indexes.
- Treat as “suggestions”, not truth.
- You still validate with workload and plan.

Use Case 3 — Index Fragmentation “Myth” (Rebuild vs Reorganize decision)

Scenario

A DBA schedules weekly rebuild for *all indexes*.

This increases:

- log generation
- maintenance time
- blocking (if not online)

...but performance doesn't improve.

Goal

Teach DBAs to take **data-driven decision**.

Step A: Check fragmentation

```
SELECT
    ips.object_id,
    OBJECT_NAME(ips.object_id) AS table_name,
    ips.index_id,
    ips.avg_fragmentation_in_percent,
    ips.page_count
FROM sys.dm_db_index_physical_stats(DB_ID(), NULL, NULL, NULL, 'SAMPLED') ips
WHERE ips.page_count > 1000
ORDER BY ips.avg_fragmentation_in_percent DESC;
```

Line-by-line explanation

- `sys.dm_db_index_physical_stats(...)`
Returns fragmentation info for indexes.
- `DB_ID()`
Current database ID.
 - `NULL` parameters
Means “all tables, all indexes”.
 - `'SAMPLED'`

Faster than FULL for large DBs; good for daily checks.

- WHERE ips.page_count > 1000

Ignore tiny indexes—fragmentation there is meaningless.

- ORDER BY avg_fragmentation_in_percent DESC

Highest fragmentation shown first.

Step B: Decision logic (industry common)

- < 5% → do nothing
- 5% to 30% → REORGANIZE
- > 30% → REBUILD

Example code snippets

```
ALTER INDEX IX_Invoices_Status_Due ON dbo.Invoices REORGANIZE;  
GO  
ALTER INDEX IX_Invoices_Status_Due ON dbo.Invoices REBUILD WITH (FILLFACTOR =  
90);  
GO
```

Explanation

- REORGANIZE

Light maintenance, online, less log, slower than rebuild.

- REBUILD

Full recreation; heavier but resets fragmentation strongly.

- FILLFACTOR = 90

Leaves free space to reduce page splits for update-heavy indexes.

Use Case 4 — Bad Plans due to Statistics (Histogram + Multi-column Stats)

Scenario

A query filters on **correlated columns** (like country_code + risk_tier).

Without multi-column stats, optimizer assumes independence → wrong estimates → wrong join type.

Query

```
SELECT e.entity_id, e.entity_name
FROM dbo.Entities e
WHERE e.country_code = 'IN'
AND e.risk_tier = 5;
```

Why estimates go wrong

- country_code and risk_tier may be correlated.
 - Optimizer may under/overestimate rows.
-

Step A: View existing stats histogram

```
DBCC SHOW_STATISTICS ('dbo.Entities', 'IX_Entities_Country') WITH HISTOGRAM;
GO
```

Explanation

- DBCC SHOW_STATISTICS
 - Displays statistics details.
 - ('dbo.Entities', 'IX_Entities_Country')
 - Uses the stats associated with the index.
 - WITH HISTOGRAM
 - Shows distribution buckets for leading column.
-

Step B: Create multi-column statistics

```
CREATE STATISTICS ST_Entities_Country_Risk
ON dbo.Entities (country_code, risk_tier);
GO
```

Explanation

- Creates stats on both columns together.
 - Helps optimizer estimate combined selectivity better.
-

Step C: Update stats after major changes

```
UPDATE STATISTICS dbo.Entities ST_Entities_Country_Risk WITH FULLSCAN;  
GO
```

Explanation

- FULLSCAN gives most accurate histogram (heavier).
- Great for demo and for critical tables after load.

Expected plan improvement

- Better estimated rows
 - Better join choices (in larger queries)
 - More stable performance
-

Final CSC DBA Takeaways (What you want participants to learn)

- ✓ Covering indexes reduce **Key Lookups**
 - ✓ Filtered indexes are perfect for **rare but critical predicates**
 - ✓ Fragmentation is not a “weekly rebuild for everything” activity
 - ✓ Statistics = optimizer intelligence (histograms + multi-column stats)
-