

Module_4_UseCase_Wait Stats, Blocking & Concurrency (4 Hours)

Concepts

- Understanding waits vs bottlenecks
- Key waits:
 - CXPACKET
 - PAGEIOLATCH_*
 - LCK_M_*
 - SOS_SCHEDULER_YIELD
- Blocking vs deadlocks
- Deadlock graphs & root cause analysis
- Isolation levels
- Row versioning (RCSI, Snapshot)

Hands-on Labs

- Capture wait stats baselines
- Identify top waits
- Simulate blocking & deadlocks
- Analyze deadlock graphs
- Fix blocking using isolation tuning

Tools Used

- sys.dm_os_wait_stats
- sys.dm_exec_requests
- Extended Events

Create CSC Global Use Cases for DBA professionals. Also explain each line of code.

Below is a **CSC Global-grade, instructor-ready Module 4 pack for**

Wait Stats, Blocking & Concurrency (4 Hours) — written exactly for **senior DBA professionals**.

You'll get:

- Real enterprise use cases
 - Hands-on labs
 - Every line of code explained
 - Clear “what to say in class” DBA takeaways
-

DAY 3 – Wait Stats, Blocking & Concurrency

Module 4: Wait Stats, Blocking & Concurrency (4 Hours)

USE CASE 1 – “System Is Slow... But CPU Is Idle”

(Understanding Waits vs Bottlenecks)

Business Context (CSC Global)

Application team says:

“SQL Server is slow, but CPU usage is only 20%.”

DBA must prove:

- Is it CPU?
 - IO?
 - Locks?
 - Parallelism?
-

Lab 1A – Capture baseline wait stats

```
SELECT
    wait_type,
    waiting_tasks_count,
    wait_time_ms / 1000.0 AS wait_time_sec,
    signal_wait_time_ms / 1000.0 AS signal_wait_sec
FROM sys.dm_os_wait_stats
ORDER BY wait_time_ms DESC;
```

Line-by-line explanation

- `sys.dm_os_wait_stats`
Cumulative waits since last restart.
- `wait_type`
What SQL Server was waiting for.
- `waiting_tasks_count`
How many times this wait occurred.
- `wait_time_ms`
Total wait time.
- `signal_wait_time_ms`
Time waiting **after** getting resource (CPU scheduling delay).

Instructor talking point

“Waits tell us *why* SQL Server was idle.”

Key DBA Learning

🔑 High waits ≠ problem by default

You must interpret *which* waits matter.

Filter out benign waits (very important)

```
SELECT
    wait_type,
    wait_time_ms / 1000.0 AS wait_time_sec
FROM sys.dm_os_wait_stats
WHERE wait_type NOT IN
(
    'CLR_SEMAPHORE', 'LAZYWRITER_SLEEP', 'RESOURCE_QUEUE',
    'SLEEP_TASK', 'SLEEP_SYSTEMTASK', 'WAITFOR',
    'HADR_FILESTREAM_IOMGR_IOCOMPLETION', 'CHECKPOINT_QUEUE',
    'XE_TIMER_EVENT', 'XE_DISPATCHER_WAIT'
)
ORDER BY wait_time_ms DESC;
```

Explanation

- Removes **idle/background waits**
 - Focuses on **actionable waits**
-

USE CASE 2 – PAGEIOLATCH_* (IO Bottleneck)

Business Context

Compliance reports slow after data growth.

Lab 2A – Identify IO waits

```
SELECT
    wait_type,
    wait_time_ms / 1000.0 AS wait_time_sec,
    waiting_tasks_count
FROM sys.dm_os_wait_stats
WHERE wait_type LIKE 'PAGEIOLATCH%'
ORDER BY wait_time_ms DESC;
```

Explanation

- PAGEIOLATCH_*
- Waiting for data pages from disk
- Indicates **storage or memory pressure**

What to say in class

“PAGEIOLATCH = SQL waiting for disk, not CPU.”

Key DBA Learning

 **Fix IO waits with indexing, memory, or storage — not CPU tuning**

USE CASE 3 – SOS_SCHEDULER_YIELD (CPU Pressure)

Business Context

CPU spikes during month-end processing.

Lab 3A – Identify CPU pressure

```
SELECT
    wait_type,
    wait_time_ms / 1000.0 AS wait_time_sec,
    waiting_tasks_count
FROM sys.dm_os_wait_stats
WHERE wait_type = 'SOS_SCHEDULER_YIELD';
```

Explanation

- SQL workers voluntarily yield CPU
 - Indicates **CPU saturation**
-

Check scheduler queues

```
SELECT
    scheduler_id,
    runnable_tasks_count,
    active_workers_count
FROM sys.dm_osSchedulers
WHERE status = 'VISIBLE ONLINE'
ORDER BY runnable_tasks_count DESC;
```

Explanation

- runnable_tasks_count > 0
 - Threads waiting for CPU
-

Key DBA Learning

 **High SOS_SCHEDULER_YIELD = CPU bound workload**

USE CASE 4 – CXPACKET / CXCONSUMER (Parallelism Issues)

Business Context

Queries are parallel but overall system slows down.

Lab 4A – Identify parallelism waits

```
SELECT
    wait_type,
    wait_time_ms / 1000.0 AS wait_time_sec
FROM sys.dm_os_wait_stats
WHERE wait_type IN ('CXPACKET', 'CXCONSUMER');
```

Explanation

- CXPACKET
 - Threads waiting on each other
- CXCONSUMER
 - Consumer threads waiting

Instructor talking point

“Parallelism is not free speed — it trades CPU for time.”

Key DBA Learning

 **Tune parallelism via better plans, not blindly lowering MAXDOP**

USE CASE 5 – Blocking (LCK_M_* waits)

Business Context

Users complain:

“Application freezes randomly.”

Lab 5A – Identify blocking sessions

```
SELECT
    r.session_id,
    r.blocking_session_id,
    r.status,
    r.wait_type,
    r.wait_time,
    r.command
FROM sys.dm_exec_requests r
WHERE r.blocking_session_id <> 0;
```

Explanation

- blocking_session_id
 - Session holding the lock
- LCK_M_*
 - Lock waits (Shared, Exclusive, Update)

Get blocking SQL text

```
SELECT
    r.session_id,
    t.text
FROM sys.dm_exec_requests r
CROSS APPLY sys.dm_exec_sql_text(r.sql_handle) t
WHERE r.session_id IN (SELECT blocking_session_id
                       FROM sys.dm_exec_requests
                       WHERE blocking_session_id <> 0);
```

Explanation

- CROSS APPLY
 - Fetches SQL text for each blocking request

Key DBA Learning

 **Blocking is not bad locking — it's long transactions**

USE CASE 6 – Deadlocks (LCK_M_* with victim)

Business Context

Application throws:

“Transaction (Process ID ...) was deadlocked.”

Lab 6A – Capture deadlocks using Extended Events

```
CREATE EVENT SESSION Deadlock_Monitor
ON SERVER
ADD EVENT sqlserver.xml_deadlock_report
ADD TARGET package0.ring_buffer;
GO

ALTER EVENT SESSION Deadlock_Monitor
ON SERVER STATE = START;
```

Explanation

- xml_deadlock_report
 - Captures deadlock graph
 - ring_buffer
 - Stores recent events in memory
-

Read deadlock graph

```
SELECT
    CAST(event_data AS XML) AS deadlock_graph
FROM sys.fn_xe_file_target_read_file
(
    NULL, NULL, NULL, NULL
);
```

Instructor talking point

“Deadlocks are logic problems, not SQL Server bugs.”

Key DBA Learning

 **Deadlocks require query or transaction design fixes**

USE CASE 7 – Isolation Levels & Blocking Reduction

Business Context

Reports block OLTP inserts.

Lab 7A – Demonstrate blocking (Session 1)

```
BEGIN TRAN;
UPDATE dbo.Entities
SET client_name = client_name
WHERE entity_id = 1;
-- Do not commit yet
```

Session 2 – Blocked read

```
SELECT *
FROM dbo.Entities
WHERE entity_id = 1;
```

Fix using READ COMMITTED SNAPSHOT (RCSI)

```
ALTER DATABASE CURRENT
SET READ_COMMITTED_SNAPSHOT ON;
```

Explanation

- Readers use **row versioning**
 - Writers no longer block readers
-

Key DBA Learning

🔑 RCSI fixes read/write blocking without code changes

USE CASE 8 – Snapshot Isolation

```
ALTER DATABASE CURRENT SET ALLOW_SNAPSHOT_ISOLATION ON;
```

Explanation

- Enables full snapshot isolation
 - Readers see transactionally consistent view
-

FINAL CSC DBA TAKEAWAYS (Module 4)

After this module, DBAs must know:

- 1 Wait stats explain **why SQL Server is slow**
 - 2 PAGEIOLATCH = IO, SOS_SCHEDULER_YIELD = CPU
 - 3 CXPACKET ≠ always bad parallelism
 - 4 Blocking = long transactions, not locks
 - 5 Deadlocks are application logic issues
 - 6 Isolation tuning solves concurrency pain elegantly
-

Interview-Ready One-Liners

- “Wait stats tell me *why* SQL Server waited.”
 - “Blocking is a symptom, not the disease.”
 - “Deadlocks are circular waits — not resource shortages.”
 - “RCSI is the safest concurrency improvement.”
-

```
/*=====
=====
CSC Global - Full Concurrency Demo Pack
Topic: Blocking → Deadlock → RCSI Fix + Deadlock Graph
Walkthrough
Audience: Senior DBAs
```

Prereqs:

- Run in a LAB database (NOT production)
- Two query windows (Session A and Session B) in SSMS
- Permissions:
 - * ALTER DATABASE (for RCSI)
 - * CREATE EVENT SESSION (for deadlock capture) - typically sysadmin

Uses your existing tables:

dbo.Entities, dbo.ComplianceFilings

Safety:

- Uses small, targeted updates (entity_id 1 and 2)

```

=====
=====*/



-----
-----  

-- PART 0) Setup (Run ONCE by Instructor)
-----  

-----  

USE CSC_DBA_Demo; -- change if needed  

GO  

  

-- Make sure XACT_ABORT is ON for safety in labs (auto-
rollback on runtime errors)
SET XACT_ABORT ON;
GO  

  

-- (Optional) Confirm rows exist
SELECT TOP (5) entity_id, client_name, country_code FROM
dbo.Entities ORDER BY entity_id;
GO  

  

-----  

-----  

-- PART 1) BLOCKING DEMO (Read blocked by Write) - Two
sessions
-- Goal: Show blocking chain and how to observe it using DMVs.
-----  

-----  

/*-----  

SESSION A (Window 1)
-----*/
-- Run this in Session A:  

BEGIN TRAN;  

  

-- Take an X lock on a single row in Entities
UPDATE dbo.Entities
SET client_name = client_name -- no-op update still takes
locks
WHERE entity_id = 1;  

  

-- Keep transaction open to hold locks (DO NOT COMMIT YET)
-- Instructor: leave this window as-is for 30-60 seconds.  

  

/*-----  

SESSION B (Window 2)
-----*/
-- Run this in Session B (it will block under READ COMMITTED):
SET LOCK_TIMEOUT 10000; -- 10 seconds so it doesn't hang
forever in class

```

```

SELECT entity_id, client_name, country_code
FROM dbo.Entities
WHERE entity_id = 1;
-- Expected: waits / times out while Session A holds the X
lock

/*
-----*
OBSERVE (Window 2 or 3)
-----*/
-- While Session B is blocked, run:
SELECT
    r.session_id,
    r.blocking_session_id,
    r.status,
    r.command,
    r.wait_type,
    r.wait_time,
    r.last_wait_type,
    r.cpu_time,
    r.total_elapsed_time
FROM sys.dm_exec_requests r
WHERE r.session_id IN (@@SPID) -- in this window
    OR r.blocking_session_id <> 0
ORDER BY r.total_elapsed_time DESC;

-- Show lock details (who holds what)
SELECT
    tl.request_session_id AS session_id,
    tl.resource_type,
    tl.resource_database_id,
    tl.resource_associated_entity_id,
    tl.request_mode,
    tl.request_status
FROM sys.dm_tran_locks tl
WHERE tl.resource_database_id = DB_ID()
ORDER BY tl.request_session_id, tl.resource_type;

-- Get SQL text for blocked & blocker sessions
;WITH blockers AS
(
    SELECT DISTINCT blocking_session_id AS sid
    FROM sys.dm_exec_requests
    WHERE blocking_session_id <> 0
),
blocked AS
(
    SELECT DISTINCT session_id AS sid
    FROM sys.dm_exec_requests
    WHERE blocking_session_id <> 0
)

```

```

)
SELECT
    s.sid AS session_id,
    t.text
FROM (SELECT sid FROM blockers UNION ALL SELECT sid FROM
blocked) s
JOIN sys.dm_exec_requests r ON r.session_id = s.sid
CROSS APPLY sys.dm_exec_sql_text(r.sql_handle) t;

-- Cleanup for blocking demo:
-- SESSION A: COMMIT to release locks
-- (Run in Session A)
COMMIT;
GO

-----
-----
-- PART 2) DEADLOCK DEMO - Two sessions + Extended Events
capture
-- Goal: Create a real deadlock and capture the deadlock
graph.
-----
-----

/*
-----
2.1 Create an Extended Events session to capture deadlocks
(Run ONCE)
-----
-----*/
IF EXISTS (SELECT 1 FROM sys.server_event_sessions WHERE name
= 'Deadlock_Monitor_CSC')
BEGIN
    DROP EVENT SESSION Deadlock_Monitor_CSC ON SERVER;
END
GO

CREATE EVENT SESSION Deadlock_Monitor_CSC
ON SERVER
ADD EVENT sqlserver.xml_deadlock_report
ADD TARGET package0.ring_buffer
WITH (MAX_MEMORY = 4096 KB, EVENT_RETENTION_MODE =
ALLOW_SINGLE_EVENT_LOSS);
GO

ALTER EVENT SESSION Deadlock_Monitor_CSC ON SERVER STATE =
START;
GO

```

```

/*
-----
----- 2.2 Deadlock scenario
Strategy:
- Session A locks Entities(1) then tries Entities(2)
- Session B locks Entities(2) then tries Entities(1)
This creates a cycle: A waits for B, B waits for A.
-----
-----*/
/*-----
SESSION A (Window 1)
-----*/
-- Run in Session A:
SET DEADLOCK_PRIORITY NORMAL;
BEGIN TRAN;

UPDATE dbo.Entities
SET client_name = client_name
WHERE entity_id = 1;

-- Pause so Session B can lock entity_id = 2
WAITFOR DELAY '00:00:05';

-- Now try to lock the row Session B has
UPDATE dbo.Entities
SET client_name = client_name
WHERE entity_id = 2;

COMMIT; -- likely won't reach if chosen as deadlock victim
GO

/*
-----  SESSION B (Window 2)
-----*/
-- Run in Session B quickly after Session A starts:
SET DEADLOCK_PRIORITY NORMAL;
BEGIN TRAN;

UPDATE dbo.Entities
SET client_name = client_name
WHERE entity_id = 2;

-- Pause so Session A is holding entity_id = 1
WAITFOR DELAY '00:00:05';

-- Now try to lock the row Session A has
UPDATE dbo.Entities
SET client_name = client_name
WHERE entity_id = 1;

```

```

COMMIT; -- one session will fail with deadlock error 1205
GO

/*
-----
Expected Output:
One session gets:
Msg 1205, Level 13, State ...
Transaction (Process ID ...) was deadlocked on lock
resources ...
and has been chosen as the deadlock victim.
-----
*/
-----


-- PART 3) READ THE DEADLOCK GRAPH (Instructor)
-- Goal: Pull XML from ring_buffer and explain it.
-----


;WITH x AS
(
    SELECT CAST(t.target_data AS XML) AS target_xml
    FROM sys.dm_xe_sessions s
    JOIN sys.dm_xe_session_targets t
        ON s.address = t.event_session_address
    WHERE s.name = 'Deadlock_Monitor_CSC'
        AND t.target_name = 'ring_buffer'
),
d AS
(
    SELECT
        n.value('(event/@timestamp)[1]', 'datetime2(3)') AS [utc_timestamp],
        n.query('.') AS event_xml
    FROM x
    CROSS APPLY
        x.target_xml.nodes('//RingBufferTarget/event[@name="xml_deadlock_report"]') AS q(n)
)
SELECT TOP (5)
    [utc_timestamp],
    event_xml.query('(event/data/value/deadlock)[1]') AS deadlock_graph_xml
FROM d
ORDER BY [utc_timestamp] DESC;
GO

-- Stop session after demo (optional)

```

```
ALTER EVENT SESSION Deadlock_Monitor_CSC ON SERVER STATE =  
STOP;  
GO
```

Deadlock graph walkthrough with annotations (how to teach it)

When you open the XML (from deadlock_graph_xml) in SSMS, you'll see this structure:

1) <deadlock>

Top-level container: "Here is the cycle."

2) <victim-list>

Look for something like:

```
<victim-list>  
  <victimProcess id="process123" />  
</victim-list>
```

Meaning: SQL Server chose this process to kill to break the cycle.

DBA key learning: Deadlock is solved by **victim selection**, not by "waiting longer."

3) <process-list>

You'll see 2 (or more) <process> nodes (one per session). Key attributes to point out:

- id="process123"

Internal id referenced by victim-list.

- spid="57"

The SQL Server session id (SPID).

- isolationlevel="read committed (2)" (or similar)

Helps you connect deadlocks to isolation behavior.

- waitresource="KEY: ..." or OBJECT: ...

What it is waiting on.

- <inputbuf>

The exact statement text that was executing at deadlock time.

Teaching script:

“Process-list answers: *Who, what SQL, what isolation, what they waited for.*”

4) <resource-list>

This is the most important for root cause. You’ll see resources like:

- <keylock> (common) → row/key-level locks
- <pagelock> → page locks (hot pages, escalation)
- <objectlock> → table-level locks

Inside each resource:

- <owner-list> → who owns the lock
- <waiter-list> → who is waiting for it

Example idea (simplified):

```
<keylock ...>
  <owner-list>
    <owner id="processA" mode="X"/>
  </owner-list>
  <waiter-list>
    <waiter id="processB" mode="X" requestType="wait"/>
  </waiter-list>
</keylock>
```

Teaching script:

“Resource-list answers: *What was locked, who held it, who waited, and what mode.*”

5) Lock modes you’ll commonly see

- S = Shared (read)
- X = Exclusive (write)
- U = Update (common in some patterns)
- IS/IX = intent locks (higher-level coordination)

DBA key learning: Deadlocks happen because two sessions acquire locks in **different orders**.

RCSI Fix Demo (Read/Write blocking reduction)

Deadlocks are usually **write/write** ordering problems; **RCSI primarily fixes read/write blocking** (reports vs OLTP).

So after the deadlock demo, show a *very practical CSC scenario*:

“Morning report reads millions of rows and blocks inserts/updates.”

RCSI setup (run once, requires exclusive access)

```
-- Run when no one else is connected to the DB (training lab  
is perfect)
```

```
ALTER DATABASE CURRENT SET READ_COMMITTED_SNAPSHOT ON WITH  
ROLLBACK IMMEDIATE;  
GO
```

Re-run the blocking demo to show improvement

Session A (Window 1)

```
BEGIN TRAN;  
UPDATE dbo.Entities  
SET client_name = client_name  
WHERE entity_id = 1;  
-- keep open
```

Session B (Window 2)

```
-- Under RCSI, this SELECT should NOT be blocked
```

```
SELECT entity_id, client_name, country_code  
FROM dbo.Entities  
WHERE entity_id = 1;
```

What changed?

- Under normal Read Committed: reader waits on writer lock.
- Under RCSI: reader gets a **versioned row** from tempdb version store.

DBA key learning (say this clearly):

- RCSI is a **safe concurrency upgrade** for OLTP + reporting mixed workloads.
- It reduces “app freeze” complaints caused by reader/writer blocking.
- It increases tempdb usage (version store), so monitor tempdb.

Production-grade takeaways for CSC DBAs

Blocking

- Usually caused by **long transactions**, not “bad locks.”
- Fix by: shorten transactions, index to reduce lock footprint, tune isolation.

Deadlocks

- Usually caused by **inconsistent object access order** across code paths.

- Fix by: consistent ordering, proper indexing, reduce lock duration, retry logic.

RCSI

- Best for reducing **read/write blocking** (reports vs OLTP).
 - Not a magic deadlock cure for write/write cycles.
-