

- Scenario (what CSC sees in production)
- Bad plan pattern (what DBAs typically observe)
- Code to reproduce (bad)
- How to read Estimated vs Actual plan (what to look for)
- Fix options (stats / rewrite / index / Query Store plan forcing)
- Improved “good plan” approach
- Plus sys.dm\_exec\_query\_stats hooks to measure impact

## Use Case 1 — Parameter Sniffing (OLTP Proc: “Entity Filings by Status”)

### Scenario

CSC compliance portal has a stored procedure:

- Sometimes called for a **rare status** (e.g., “Overdue” = 0.1%)
- Sometimes for a **common status** (e.g., “Pending” = 60%)

First execution compiles plan for one value → reused for others → **random slowness**.

### BAD code (creates sniffing behavior)

```

CREATE OR ALTER PROCEDURE dbo.usp_GetFilingsByStatus
    @Status VARCHAR(20)
AS
BEGIN
    SET NOCOUNT ON;

    SELECT f.filing_id, f.entity_id, f.due_date, f.filing_status
    FROM dbo.ComplianceFilings f
    WHERE f.filing_status = @Status
    ORDER BY f.due_date DESC;
END;
GO

-- Bad pattern: first execution compiles plan for rare/common value and
-- reuses it
EXEC dbo.usp_GetFilingsByStatus @Status = 'Overdue';
EXEC dbo.usp_GetFilingsByStatus @Status = 'Pending';

```

### BAD plan pattern (what you’ll see)

- Plan compiled for **rare value** might choose **Index Seek + Nested Loops**
- When reused for **common value**, it becomes expensive:
  - lots of lookups
  - high CPU/reads
- Or vice versa: compiled for common value → **Scan/Hash** reused for rare value → slow.

## What to check in Estimated vs Actual plan

- **Estimated rows vs Actual rows** at the main access operator
- Join choice, key lookups, memory grant changes

## Fix options (in order)

### Fix A (quick + safe for skew):

#### OPTION (RECOMPILE)

```
CREATE OR ALTER PROCEDURE dbo.usp_GetFilingsByStatus
    @Status VARCHAR(20)
AS
BEGIN
    SET NOCOUNT ON;

    SELECT f.filing_id, f.entity_id, f.due_date, f.filing_status
    FROM dbo.ComplianceFilings f
    WHERE f.filing_status = @Status
    ORDER BY f.due_date DESC
    OPTION (RECOMPILE);
END;
GO
```

Best when execution frequency is moderate and skew is high.

### Fix B (stabilize with Query Store plan forcing)

Use Query Store after identifying the good plan (see Use Case 5).

### Fix C (index to support both patterns)

If you frequently filter + sort:

```
CREATE INDEX IX_Filings_Status_DueDate
ON dbo.ComplianceFilings (filing_status, due_date DESC)
```

```
INCLUDE (entity_id);
GO
```

---

## Use Case 2 — Bad Estimates from Stale Statistics (Join flips wrong)

### Scenario

After data load/month-end, queries that join **Entities** → **Filings** become slow.

No code change. Root cause: **stale stats** → wrong row estimates → wrong join choice.

### BAD query

```
SELECT e.entity_id, e.entity_name, f.filing_id, f.due_date
FROM dbo.Entities e
JOIN dbo.ComplianceFilings f
  ON f.entity_id = e.entity_id
WHERE e.country_code = 'IN'
  AND f.due_date < DATEADD(day, 30, GETDATE());
```

### BAD plan pattern

- Huge mismatch between **Estimated vs Actual rows**
- Optimizer may pick:
  - Nested Loops + repeated lookups (bad if result set large)
  - Or Hash Join with huge memory grant/spills (if estimate is wrong)

### Fix A — Update stats (fastest “proof” fix)

```
UPDATE STATISTICS dbo.Entities WITH FULLSCAN;
UPDATE STATISTICS dbo.ComplianceFilings WITH FULLSCAN;
GO
```

### Fix B — Add/adjust index to match predicates

```
CREATE INDEX IX_Entities_Country
ON dbo.Entities (country_code)
INCLUDE (entity_name);
```

```
GO
```

```
CREATE INDEX IX_Filings_Entity_DueDate  
ON dbo.ComplianceFilings (entity_id, due_date)  
INCLUDE (filing_id);  
GO
```

## What to teach DBAs

- If estimates are wrong → plans are unstable
  - Stats freshness is *often* the first lever before rewriting everything
- 

## Use Case 3 — Scan vs Seek due to Non-SARGable Predicate (CPU + logical reads spike)

### Scenario

A support query is written like this:

- It “looks right”
- But causes **table/index scan** and heavy logical reads

### BAD query (non-sargable)

```
SELECT f.filing_id, f.entity_id, f.due_date  
FROM dbo.ComplianceFilings f  
WHERE CONVERT(date, f.due_date) = CONVERT(date, GETDATE());
```

### BAD plan pattern

- **Index Scan** even if you have an index on due\_date
- High logical\_reads

### GOOD rewrite (sargable range predicate)

```

DECLARE @d date = CONVERT(date, GETDATE()) ;

SELECT f.filing_id, f.entity_id, f.due_date
FROM dbo.ComplianceFilings f
WHERE f.due_date >= @d
    AND f.due_date < DATEADD(day, 1, @d) ;

```

## Supporting index (if needed)

```

CREATE INDEX IX_Filings_DueDate
ON dbo.ComplianceFilings (due_date)
INCLUDE (entity_id, filing_id);
GO

```

## What to check in plans

- BAD: Scan + residual predicate
  - GOOD: Seek + predicate pushdown + lower logical reads
- 

## Use Case 4 — Join Order Problem (filter applied too late → huge work)

### Scenario

A report joins big tables then filters afterwards.

Result: optimizer may process far more rows than needed.

### BAD query (filter applied late / shape not helping optimizer)

```

SELECT e.entity_name, f.filing_id, f.due_date
FROM dbo.Entities e
JOIN dbo.ComplianceFilings f
    ON f.entity_id = e.entity_id
WHERE f.filing_status IN ('Pending', 'Overdue')
    AND e.entity_name LIKE '%TECH%';

```

### BAD plan pattern

- Large intermediate row sets
- Hash join + large memory grant or spills
- High elapsed time

## GOOD rewrite (reduce rows early)

```
;WITH FilteredEntities AS (
    SELECT entity_id, entity_name
    FROM dbo.Entities
    WHERE entity_name LIKE '%TECH%'
),
FilteredFilings AS (
    SELECT filing_id, entity_id, due_date, filing_status
    FROM dbo.ComplianceFilings
    WHERE filing_status IN ('Pending','Overdue')
)
SELECT e.entity_name, f.filing_id, f.due_date
FROM FilteredEntities e
JOIN FilteredFilings f
ON f.entity_id = e.entity_id;
```

## Index support (usually required)

```
CREATE INDEX IX_Entities_Name
ON dbo.Entities (entity_name)
INCLUDE (entity_id);
GO

CREATE INDEX IX_Filings_Status_Entity
ON dbo.ComplianceFilings (filing_status, entity_id)
INCLUDE (due_date, filing_id);
GO
```

## Teaching point

- “SQL is declarative” but **query shape influences cardinality + join order**
- Help optimizer by filtering early + indexing filter columns

## Use Case 5 — Plan Regression after Deployment (Use Query Store to force “good plan”)

### Scenario

After a deployment:

- Query becomes slow
- Plan changed
- DBAs need **fast rollback without code change**

## Step 1 — Ensure Query Store is ON

```
ALTER DATABASE CURRENT SET QUERY_STORE = ON;
ALTER DATABASE CURRENT SET QUERY_STORE (OPERATION_MODE = READ_WRITE);
GO
```

## Step 2 — Identify top resource queries (Query Store view)

```
SELECT TOP (20)
    qsqt.query_sql_text,
    qsp.plan_id,
    rs.avg_duration,
    rs.avg_cpu_time,
    rs.count_executions
FROM sys.query_store_runtime_stats rs
JOIN sys.query_store_plan qsp
    ON rs.plan_id = qsp.plan_id
JOIN sys.query_store_query qsq
    ON qsp.query_id = qsq.query_id
JOIN sys.query_store_query_text qsqt
    ON qsq.query_text_id = qsqt.query_text_id
ORDER BY rs.avg_duration DESC;
```

## Step 3 — Force the known good plan

```
EXEC sys.sp_query_store_force_plan @query_id = 123, @plan_id = 456;
```

## Step 4 — Validate improvement

- Compare runtime stats for forced plan
- Verify no new side effects
- If needed:

```
EXEC sys.sp_query_store_unforce_plan @query_id = 123, @plan_id = 456;
```

- 
- **Hints** are “hard force” and often permanent debt
  - **Query Store forcing** is safer, reversible, and auditable

# Measuring “Bad vs Good” with sys.dm\_exec\_query\_stats (training-friendly)

Use this to show impact before/after (CPU, reads, elapsed):

```
SELECT TOP (20)
    qs.total_worker_time / 1000 AS total_cpu_ms,
    qs.total_elapsed_time / 1000 AS total_elapsed_ms,
    qs.total_logical_reads,
    qs.execution_count,
    (qs.total_worker_time / NULLIF(qs.execution_count,0)) / 1000 AS
    avg_cpu_ms,
    SUBSTRING(st.text, (qs.statement_start_offset/2)+1,
        ((CASE qs.statement_end_offset WHEN -1 THEN DATALENGTH(st.text)
        ELSE qs.statement_end_offset END - qs.statement_start_offset)/2)+1)
AS statement_text
FROM sys.dm_exec_query_stats qs
CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) st
ORDER BY qs.total_worker_time DESC;
```

---